

C++ Training

Introduction and basics

Olivier Parcollet

Center for Computational Quantum Physics (CCQ)

Flatiron Institute, Simons Foundation

New York



Outline

Thursday Oct. 24	9:00-9:15am	Introduction : Principles
	9:15-10:00am	Basics
	10:00-10:15am	Break
	10:15-11:00am	Standard Library: Containers & Algorithms
	11:00-11:30am	Lambda Functions
	11:30-12:00pm	STL Algorithms II
	12:00-1:00pm	Lunch
	1:00-1:40pm	Editing & Static Analysis
	1:40-2:10pm	Git Basics
	2:10-2:30pm	Break
	2:30-3:00pm	Github, Pull requests and Code Review
	3:00-3:20pm	Managing Config Files
	3:20-3:40pm	Continuous Integration
	3:40-4:25pm	CMake and App4Triqs
	4:25-5:10pm	Testing

Friday Oct. 25	9:00-9:20am	Error handling
	9:20-9:50am	Move semantics
	9:50-10:00am	Break
	10:00-11:00am	Concepts
	11:00-12:00pm	TRIQS library components
	12:00-1:00pm	Lunch
	1:00-2:30pm	C++ HandsOn
	2:30-3:30pm	Debugging
	3:30-3:45pm	Break
	3:45-4:15pm	Interfacing C++ and Python
	4:15-4:30pm	Documenting
	4:30-5:00pm	Profiling & Optimizing

Targeted audience

- Upgrade C++ knowledge for e.g.
 - C programmers (use higher level abstraction).
 - Python programmers (I will show some similarity).
- I assume some familiarity with the basic (?) syntax of C++
- Do not hesitate to ask questions !
- NB : Most of the material is pure C++, some TRIQS example
- Tomorrow : a little overview of TRIQS library components.

Level of C++ expertise

1. User level
to write code using modern C++ libraries and techniques e.g. TRIQS
[this session]
2. Library writer level
Not the purpose of this session (a few notions only).

Why C++ ?

- Why not Python, Rust, D, Swift, Julia, ?
- Industry standard. ISO.
- **Stability** but also **evolution** (“modern” C++).
- Performance.
 - **Zero cost/overhead abstractions.**
= write more expressive code, without performance penalty.
- A language designed to build libraries.
- Largely used in HPC.

What is zero cost abstraction ?

- What is *simple* should be coded *simply*
- High level and yet fast.
- Very important for readability, long term maintenance, code review.
- Common wrong idea : compact, simple, readable code is slow.
- We want simplicity (**abstraction**), without any performance penalty (at **zero cost**).
- Generic programming is essential to achieve this.
C++17, C++20 make it a lot easier.

Simple example

- Let us look at a tiny piece of code.
- With our matrix class
(triqs::arrays, soon pulled out of TRIQS).

```
// using nda = triqs::arrays;  
  
int n = 5, p = 6;    // any number ...  
nda::matrix a(n, p); // create a matrix of dimension n x p  
-----  
// Put all elements to 0  
// Version 1  
a = 0;  
-----  
// Version 2  
for (int i = 0; i < n; ++i)  
    for (int j = 0; j < p; ++j)  
        a(i, j) = 0;
```

- Which one is better : version 1 or version 2 ? Why ?
- Good code expresses intent, not implementation details.

A puzzle for tomorrow

- A and B : two matrices $n \times n$, real valued.

A function *trace*

- We want to write

```
double r = trace (A + B);
```

- Instead of

```
double r = 0;  
for (int i = 0; i < n; ++i)  
    r += A(i, i) + B(i, i);
```

- A priori, zero cost abstraction seems impossible:
 - $A + B$ computed first, before calling `trace`.
 - Scales as N^2 while hand-written code is N

What is the cost of abstraction ?

Compilation is a transformation

- Compilers do not exactly implement what you write ...
... but a code with the same visible results
- Compilers transform code (a lot), inline functions, rewrite loops (vectorization, etc...), eliminate code with no effect.

gcc-explorer tool

11

- Give a small piece of C++, it compiles, executes and shows the assembly code.
- Online : <https://godbolt.org>
- NB : No deep knowledge of assembly required !

Demo

Compilers can surprise you ...

<https://godbolt.org/z/8HeLqh>

- A simple function

The screenshot shows the Godbolt compiler explorer interface. The source code on the left is a C++ function `calc(int n)` that calculates the sum of integers from 0 to `n` using a loop. The assembly output on the right shows the compiler's optimization. It replaces the loop with a direct calculation: `ax = n*(n-1)/2 + n`. The assembly also includes a branch to `.LBB0_1` if `n` is zero, where it returns 0.

```

11010 .LX0: .text // Intel
1 calc(int):
2     test    edi, edi           if n==0 :
3     js      .LBB0_1           goto .LBB01
4     mov     ecx, edi           n
5     lea     eax, [rdi - 1]     ax = n-1
6     imul    rax, rcx           ax = n*(n-1)
7     shr     rax               ax = ax / 2
8     add     eax, edi           ax = ax + n
9     ret                                return ax
10 .LBB0_1:
11     xor     eax, eax
12     ret                                return 0

```

- Compiler has replaced code by `return n*(n-1)/2 + n`
- Program speed sometimes hard to predict. **Measure !**

C++14, 17, 20 ??

C++ evolution


- A new standard now every 3 years :
C++98, C++11, C++14, C++17, C++20, ...
- Pushed by industry (e.g. Google, from mobile to large server).
- C++ is becoming simpler for users, for library writers.
- High-level constructs inspired by e.g. Python, Ocaml, ...
- Standard library richer.
- Backward compatibility very rigorously enforced.
 - No change of meaning of 1/2 like in Python 2 to 3 !!

Compilers


- Main compilers : `gcc (GNU)`, `llvm/clang`.
- Many others, e.g. Intel.
- Our policy : upgrade compilers often. Use latest versions.
 - Gcc release every year, llvm/clang every 6 months.
 - Yes ! Benefits >> Costs (our decision for TRIQS).
 - Technical solution (e.g. singularity).

Forget old books !


- How to write good C++ has changed.
- Training materials must be updated.



Book by P. Gottschling



Books by B. Stroustrup




- Please, do not take random learning material on the internet !

Reference ?

17

- On the web : <http://en.cppreference.com/w/>

- Reference book



Where to get information ?

- ISO C++ committee. E.g. check for information <http://isocpp.org>
- Consensus style. Backward compatibility.
- Slow, but very high quality

The screenshot shows the homepage of the C++ Standard website. The header includes navigation links: Get Started!, Tour, Core Guidelines, Super-FAQ, Standardization, and About. A search bar is located in the top right corner. The main content area is titled "News, Status & Discussion about Standard C++".

Follow All Posts
The home of Standard C++ on the web – news, status and discussion about the C++ standard on all compilers and platforms.

Recent Highlights

- [Slides of the 28th of June 2018 BeC++ Meeting—Marc Gregoire](#)
By Marc Gregoire | Jul 6, 2018 09:34 AM
- [CppCon: New promo video, and Early Bird registration ends this weekend](#)
By Blog Staff | Jul 6, 2018 08:34 AM
- [CppCon keynote: Patterns and Techniques Used in the Houdini 3D Graphics Application \(Mark Elendt\)](#)
By Blog Staff | Jul 5, 2018 01:42 PM
- [C++ on Sea: Call for Speakers](#)
By Felix Petriconi | Jul 5, 2018 07:33 AM
- [CopperSpice: Special Member Functions](#)
By Ansel Sermersheim | Jul 5, 2018 07:18 AM

Articles & Books

- [CppCon 2017: Agent based class design—Odin Holmes](#)
By Adrien Hamelin | Jul 4, 2018 01:09 PM
- [Parallel STL And Filesystem: Files Word Count Example—Bartłomiej Filipek](#)
By Adrien Hamelin | Jul 2, 2018 12:39 PM
- [Functions of Variants are Covariant—Alfredo Correa](#)
By Adrien Hamelin | Jul 2, 2018 12:28 PM
- [The surprisingly high cost of static-lifetime constructors—Arthur O'Dwyer](#)
By Adrien Hamelin | Jun 27, 2018 12:13 PM

Selected Recent C++ Questions
 Tagged with C++, C++11, C++14, C++1z, or C++17

- [Are there certain keywords that should not be "#defined" by me?](#)
By Markali | Jul 7, 2018 11:56 AM
- [Why does `emplace_back\("Hello"\)` call `strlen`?](#)
By Cornstalks | Jul 7, 2018 5:53 AM
- [Specifying `-std` when compiling both C and C++ code](#)
By Siguza | Jul 7, 2018 2:03 AM
- [Iterating a vector to second to last element with index vs iterator](#)
By chew socks | Jul 6, 2018 10:26 PM
- [Is it possible in modern C++ to pass a string literal as a parameter to a C++ template?](#)
By Bill Moore | Jul 5, 2018 5:47 PM
- [Can `std::is_invocable` be emulated within C++11?](#)
By Janik | Jul 5, 2018 11:36 AM

Product News

- [Announcing C++ Just My Code Stepping in Visual Studio—Marian Luparu](#)
By Marco Arena | Jul 2, 2018 01:46 AM
- [ZAPCC compiler is now available under Open Source—Ceemle](#)
By Felix Petriconi | Jun 18, 2018 05:04 AM
- [PVS-Studio 6.24 released](#)

FEATURES

- Current ISO C++ status
- Upcoming ISO C++ meetings
- Compiler conformance status

TAGS

- basics
- intermediate
- advanced
- experimental

UPCOMING EVENTS

- CppCon**
Sep 23-29, Bellevue, WA, USA
- pacific++**
Oct 18-19, Sydney, Australia
- code::dive 2018**
Nov 7-8, Wrocław, Poland
- Meeting C++**
Nov 15-17, Berlin, Germany
- C++ on Sea**
Feb 4-6, Folkestone, Kent, UK

Tweets by @isocpp

Standard C++
@isocpp

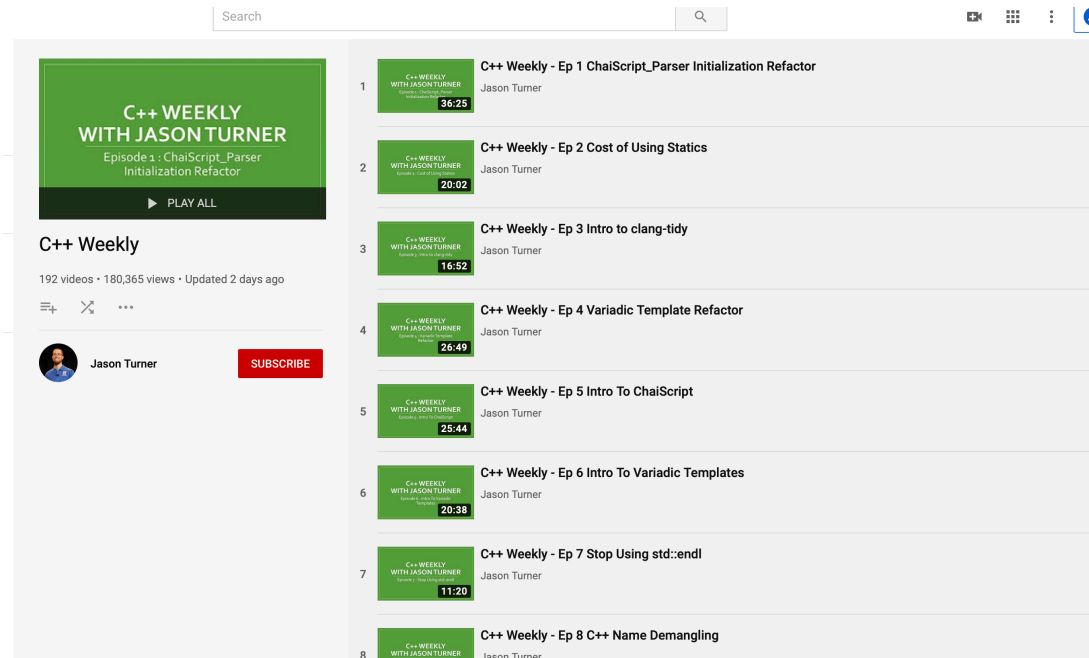
Where to get information ?

- Blog, e.g. “Fluent C++”

<https://www.fluentcpp.com>

- C++ weekly on youtube

<https://www.youtube.com/user/lefticus1>



Subset of C++

Subset of C++

- C++ is a vast, multipurpose language
- “Multi-paradigm” :
imperative functional programming,
generic programming, object orientation
- We use a subset of the language
- Discard some old features (too verbose, unsafe)
Instead use some recent features.
- Also some C++ (mandatory) style guidelines
- Which subset ? How to enforce it (compiler, tool) ?

Subset of C++ : highlights

- No inheritance (almost).
- Prefer the functional style :
 - Functions that do not change their arguments (purity)
 - “Data oriented”
 - Use good libraries and containers.
- Use generic programming and compile time decisions.
Becoming mainstream.

Coding guidelines

- How to enforce a subset of C++ ? (like “safe” part of D e.g.).
- Rules and/or recommendations for coding style
 - e.g. C++ Core Guidelines
 - <https://github.com/isocpp/CppCoreGuidelines/blob/master/CppCoreGuidelines.md>
 - <https://isocpp.github.io/CppCoreGuidelines/CppCoreGuidelines#Rr-ptr>
- Team/project specific (e.g. TRIQS).

Static analysis of code

- **Compiler**
 - Basic, but you should use -Wall -Wextra et al.
 - (Most) compilers warnings must be resolved
- **clang-tidy**
 - A tool of the clang family
 - Detects violation of guidelines, some bugs, it can even fix/modernize the code.
 - Choose which guidelines to enforce in the team
 - Some are integrated in IDE, e.g. CLion, VsCode.

Basics

A selection of topics

Value and references

- C++ manipulates two very different kind of things.
 - Regular types.
 - References or pointers.

Regular types

- Obey a set of simple rules/axioms :

- Copying make a new object

- Assignment is deep.

- Default constructible

- ...

- Simplest example : int, double.

- `std::vector<T>`
*a simple container with N
 elements of type T,
 contiguous in memory*

```
// comparison follows from copy
T a = b; assert(a==b);

// copy and assignment are the same
T a1; a1 = b; T a2 = b; assert(a1 == a2);

// copy/assignment is by value, not reference
T a = c; T b = c; a = d; assert(b==c);
```

```
int a = 3;

int b1 = a;

int b2;
b2 = a;

std::vector<int> a = {1, 2, 3};

std::vector<int> b1 = a;

std::vector<int> b2;
b2 = a;
```

- Try to make your types Regular

Regular type (2)

- Composition

```
struct A {  
    int N;  
    double x;  
    B b; // B is regular  
};
```

- A is a regular type.
- Compiler can generate good default for copy, assignment, ...

Reminder : template class in C++

```
template<typename T>  
class vector {  
  
    // some array of T and many methods...  
  
};
```

- For each T, the compiler will generate a new code

Reminder : template function in C++

```
template <typename T> T inc(T x) {  
    return x + 1;  
}
```

- A function for all type T
- inc can be called for an int, double, anything for which $x + 1$ makes sense.
- In each case, the compiler will generate a new piece of code, optimize it, etc.
- Called “generic programming”

References : reminder

```
int i = 10;  
  
int & r = i; // A reference to i  
  
r = 3; // now i AND r contains 3  
  
i = 4; // now i AND r contains 4  
  
int const & r = i; // a reference but I can not change its value.
```

- Another name of the same variable.
- References are :
 - Very cheap to construct. No copy of the object.
 - Immutable (one can not reassign them).

Reminder : ref usage

- Pass by reference

```
int f(A const & a) {  
    // use a, no copy, no modification  
}  
int f(A & a) {  
    // Can modify a, no copy.  
}
```

- See a part of a big object

```
class A {  
    std::vector<int> _data;  
    int something_else;  
  
    public:  
    std::vector<int> const & data() const { return _data; }  
};  
  
// usage  
A a;  
do_something(a.data());
```


Reminder : const

- Something that should not be changed

```
int const a = 10;  
  
int b = 23;  
auto const & c = b;
```

```
int f(A const & a) {  
    // use a, no copy, no modification  
}  
  
int f(A & a) {  
    // Can modify a, no copy.  
}
```

```
struct A {  
    int u = 10;  
    int g(int v) const {  
        // the method does not change u  
        return u+v;  
    }  
};
```

Example of non regular type

```
struct A {  
    int N;  
    double & x;  
};
```

- Issues : copy ? Use with STL, generic algorithm ?
- Try to avoid such types, if not necessary.

const correctness

const correctness

- Make everything const by default everywhere possible.
- Const correctness is painful to fix at a later stage.

Good

```
int f (A const & a) {
    // a can not be changed.
}
```

Compiles but **bad**
Does not pass code review

```
int f (A & a) {
    // if a is not changed.
    // compiles but bad style
}
```

Intent is : f modify x

```
struct A {
    int u = 10;
    int g(int v) const { return u+v};
};
```

```
struct A {
    int u = 10;
    int g(int v) { return u+v};
};
```

Pointers...

Pointers : reminder

- A pointer is the address of a variable
- Similar to references.
Differences : Syntax. Can be nullptr, can be reassigned.
Can be in an invalid state.

```
int i = 10;  
  
int *p = &i; // pointer to i  
  
*p = 3; // now i contains 3  
  
i = 4; // now *p is 4  
  
int const *p = i; //
```

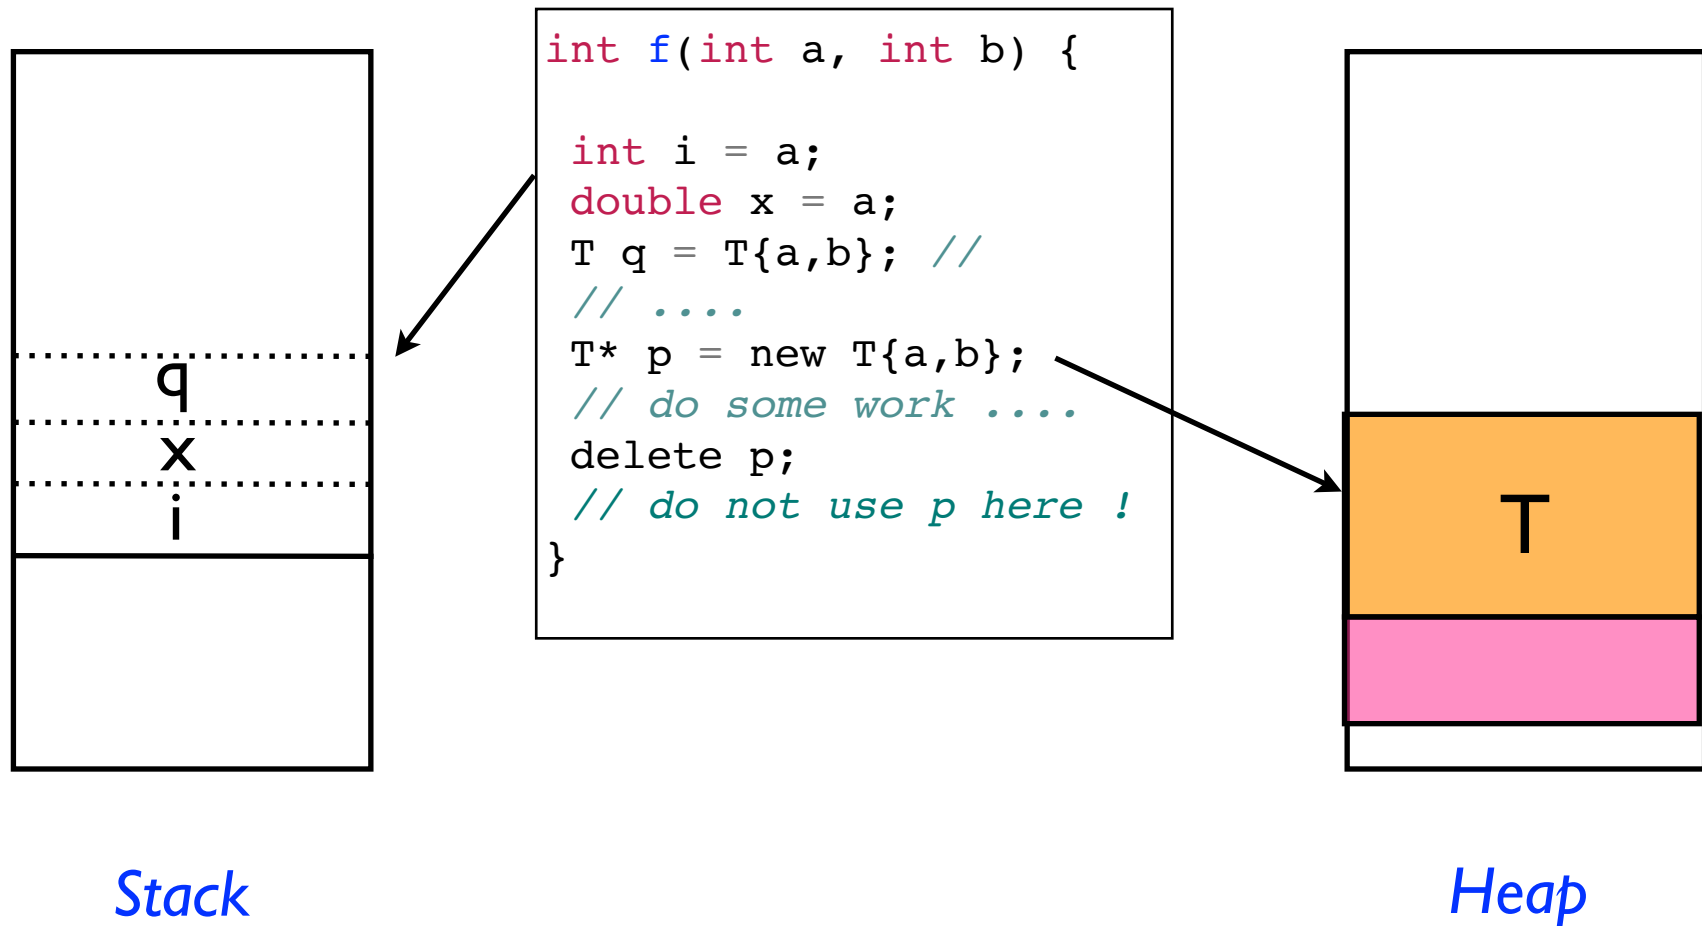
```
int i = 10;  
  
int & r = i;  
  
r = 3;  
  
i = 4;  
  
int const & r = i; //
```

Pointers in modern C++

- Pointers (or iterators) are everywhere in C / C++
- Modern C++ use much less pointers

Reminder : stack vs heap

- Stack allocation vs heap allocation




- Local variables
- Destroyed automatically at }

- Outlives } unless explicitly deleted

Classical C/C++ issue with pointers

```
T* make_big_object_on_heap(int a, int b)
{
    T* p = new T{a,b};
    return p;
}
```



- p owns the data on the heap.
Who is in charge of deleting p ?
- Two classical type of bugs in C/C++:
 - **Memory leaks** : nobody takes care of delete
 - **Dangling pointers** : use after delete. segfault ! ?
- Solution :
 - We can now detect these bugs easily (soon even at compile time !)
 - Use higher level abstractions which avoid these issues

Heap

Modern C++ : Use containers

- `std::vector`, array/matrix class, `gf` class in TRIQS, ...
- They handle the memory allocation for you
- They are regular types
- C++ style :
 - Do NOT use owning raw pointers
 - NB : Pointers are fine to observe an object, without ownership.

Returning a large object

- A function that builds and return a large object
- e.g. `std::vector`, matrix, array

```
BigObject f(int a, int b) {  
    // ....  
    return BigObject{ whatever(...) };  
}
```


- Simplest code (just return it) is also the most efficient
- Before C++11, there was a copy at return (even if compilers sometimes optimized it).
- Hence many books have obsolete patterns : return by pointer, etc..
- *Return Value Optimization (RVO)*. Cf also move semantics, day 2.

Smart pointers

- If no container is available (most likely you will not need it)


- `std::unique_ptr`

- One pointer max to each object
- `p` can not be copied.



- `std::shared_ptr`

- Multiple pointers
- With reference counting, like Python
- When `#ref = 0`, object is destroyed



Guideline

- Use **Containers** (std::vector, array/matrix class, gf class...)
- Use **smart pointers** (unique_ptr first, shared_ptr maybe ?)
- Do NOT use **owning raw pointers**
 - **Never use new/delete** (does not pass code review !)

Scopes

}

Scopes : reminder

- Delimited by { }

```
{  
    int i =9;  
    std::vector<int> v= {1,2,3,4};  
  
    // work  
} // everything is cleaned up here ...
```

```
// a function is a scope  
int f(int n) {  
  
}
```

- Each object has a constructor (at initialization) ...
- ... and a destructor (executed at }
- Guaranteed to be executed, even in case of exceptions.

```
struct A {  
    A (int i);  
    ~A();  
};
```


RAII pattern

- A common C++ idiom : Resource Acquisition Is Initialization (?!).

```
// code defined some variable  
double a;  
gf<imtime, matrix_valued> { whatever  };  
  
{  
    h5::file f("myfile.h5", 'w');  
  
    h5_write(f, a, "A");  
    h5_write(f, g, "g");  
  
} // DONE. File is closed here because f is destroyed.
```

- Similar to Python “with...”
- Releasing resource (memory, file, ...) is **automatic**

RAII : example of unique_ptr

- `std::unique_ptr`
pointer to an object, no other pointer pointing to it
- Not copy, can be moved.

```
{  
  T* p = new T{a,b};  
  //  
  // do a lot ....  
  
  delete p;  
}
```

Bad

Should not pass code review

```
{  
  std::unique_ptr<T> p = std::make_unique<T>(a,b);  
  // work ...  
} // delete is automatique
```

Ok

```
{  
  T x{a,b};  
  // work ...  
} // delete is automatique
```

Even simpler

auto

auto : automatic type deduction

- Ask the compiler to deduce the type for us

```
auto y = x;           // makes a copy of x  
auto z = f(a,b);      // same from the result of f
```

```
auto v = std::vector<int>{1, 2, 3};
```

- NB : the type is still fixed at compile time (not dynamical as in Python)

- Enforce a type ?

```
auto x = whatever;
```

VS

```
auto x = T{whatever};
```

```
T x {whatever};
```

- Declare multiple variables. *Structure binding*

```
auto [x, y] = std::make_tuple(1, 4.3);
```

Almost always auto (AAA)

- We recommend to use auto in most place

```
auto a = A { arguments};  
  
auto result = my_function(a);  
  
auto m = my_matrix{ ...};  
auto g = gf<imtime>{...};  
  
int i =0;
```

- Makes code simple and more regular
- Helps a lot in generic code. Can save from unwanted conversion.
- Sometimes, e.g. for lambda, there is no other choice (Cf Nils' talk).

Loops


A simple loop

- A simple loop in Python ...

```
v = [1,3,5,9]
s = 0
for x in v:
    s+=x
```

- ... C++ equivalent. Main difference is types.

Intuitive



```
auto v = std::vector<int> {1,3,7,9};
int s = 0;
for (auto x : v) {
    // do something ...
    s+= x;
}
```

Loops : avoid copies

```
auto v = std::vector<int> {1,3,7,9};
int s = 0;

for (auto x : v) {
    // do something !
    // may be quite complex
    s += x;
}
```

- More generally : const, not const versions

```
auto v = std::vector<BigType> { /*...*/ };

// v is unchanged, all elements are visited.
for (auto const& x : v) {
    // ...
}

// all elements visited, they can change
for (auto& x : v) {
    // ...
}
```


Simpler than what ?

- Modern C++

```
for (auto const & x : v) {  
    // do something !  
    s+= x;  
}
```

- Intent is clearer :
 - . Iterate on every elements in order
 - . v unchanged
- As or more efficient.

- Old C++

```
for (std::vector<int>::const_iterator it=  
v.begin(); it != v.end(); ++it) {  
    // do something !  
    s+= *it;  
}
```

```
for (int i = 0; i < v.size(); ++i)  
{  
    // do something !  
    s+= v[i];  
}
```

Like Python itertools

- More sophisticated iterations.

```
#include <itertools/itertools.hpp>
std::vector<int> vec;
// ...
for (auto [n, x] : itertools::enumerate(vec)) {
    // (0, x[0]), (1, x[1]), (2, x[2]), ...
}
```

```
std::vector<int> vec2, vec1;
// ...

for (auto [x, y] : itertools::zip(vec1, vec2)) {
    // (x[0], y[0]), (x[1], y[1]), (x[2], y[2]), ...
}
```

- A simple header file, pulled out of TRIQS. Apache 2 licence.
- C++20 : ranges will be part of C++ std library.

A quick look into C++20

- Implementing such things in C++ will become very simple ...
- **Coroutines**
 - Generators like in Python.
 - And much more ...

```
def enumerate(X) :  
    n=0  
    for x in X:  
        yield n, x  
        n +=1
```

Python

```
template<typename T>  
std::generator<std::pair<n, typename T::value_type>>  
  
enumerate(T const & x) {  
    int n=0;  
    for (auto const & y : x) {  
        co_yield std::pair{n,y};  
        n++;  
    }  
}
```

C++20

- NB : still use libraries by default, do NOT reimplement yourself...

Summary of basic guidelines

- Use Regular types
- Use containers. Do not use raw owning pointers
- Use const by default
- RAI : destructors clean after you.
- Use auto (AAA)
- Use expressive loops

Thank you for your attention