

Debugging

Olivier Parcollet

Center for Computational Quantum Physics (CCQ)

Flatiron Institute, Simons Foundation

New York

SIMONS FOUNDATION



Docker image for demo

```
docker pull wentzell/debugging_demo
```

Outline

- What kind of bugs ?
 - A selection of (open source) tools
 - **Examples.** Tiny pieces of codes, **show the tools in action**
 - Bug prevention.
-
- Here simple illustrations with (basic) C++
 - Many tools usable in Fortran, similar tools in most languages.
 - A few C++ specific issues.

Tests !

- How to detect bugs : **have a good test suite.**
- **Test driven development**
- It much better to kill bugs in a test than in a production code.
- **Test coverage :**
 - Are all functions tested ? For which parameters ?
 - If you do not RUN your function, it is harder to find bugs (but see static analysis, cf later).
 - There are tools to measure that (kcov, ...). Not the topic here.

Different kind of bugs

- Out of bound read/write
 - e.g. array of size N , write/read in $N+10$.
- Memory leak
- Uninitialised variables
- Undefined behaviour
- Logic error
- Infinite loop
- ...

Tools

- Static analysers

- Static analyzers : detect errors in the code without running it.
- Helps you to write better, cleaner code. Enforce subset of C++.
- Compiler itself !, clang-tidy, cppcheck, ...

- Dynamical analyzers

- Automated checks of the code while running (check bounds, etc)
 - Valgrind: a “virtual” machine, checking each memory access, that each variable is initialized. No recompilation. Fortran/C++.
 - Clang/gcc sanitizers : compilation options that “instrument” your code automatically to do these checks. (clang/gcc only).


- Debugger

- Execute program step by step : GDB, LLDB, ...

Static/Dynamical checks

- Use them preventively, on your test suite/code
- Systematic use in “Continuous Integration” (Travis, Jenkins)
- Static analysis
 - Check all code.
 - Does not require tests
 - False positives possible
- Dynamic analysis
 - Require very good test coverage
 - Require to run the tests
 - Almost no false positive

Valgrind



- <http://valgrind.org>
- **Best on Linux** (does not work well on OS X).
- Various checkers : here we use “memory” (the default one).
- Contains other tools to analyse cache misses, profiling
CallGrind, CacheGrind

Clang sanitizers



clang.llvm.org

- Dynamic analysis tools included with clang and gcc
- **Address Sanitizer** `-fsanitize=address`
detects buffer overflows, memory leaks, use-after-free,...
 - **Memory Sanitizer (clang only)** `-fsanitize=memory`
detects reads of uninitialized memory
 - **Undefined Behavior Sanitizer** `-fsanitize=undefined`
detects undefined behaviors
 - **Thread Sanitizer** `-fsanitize=thread`
detects data races
- Used to systematically detect bugs by Google, Mozilla, ...
 - Much less memory/runtime overhead than Valgrind!

Availability

- Linux
 - Clang (>3.x), gcc (>4.9)
 - All sanitizers
- Os X
 - Not all sanitisers
 - Valgrind does not work well.
- Others : ?
- Recommendation : use Linux for debugging

Examples

I-Out of bounds !

- Typical error : read/write out of the memory reserved for an object.

```
#include <vector>

int fun(int n) {
    std::vector<int> vec(100, 1); // a 1d "array" of 100 int init to 1
    vec[n + 50] = 500;           // Invalid Write
    return vec[n + 50];          // Invalid Read
}
```

- Similar : use a “dangling” or null pointer, ...
- How does it manifest itself ?


Demo

I-Out of bounds !

- Segfault : the program crashes ...
- ... or not !
- Hence it **can be undetected for a while (or never ?)**.
- We need to find **the first error**



Computer
memory



Computer
memory

Segfault: crash

*Even worse. No crash,
but corrupt memory*

OK

1-Out of bounds !

- Debugger is NOT the best tool in this case.

- The code may not crash
- We want a systematic check
- We want the first error

- Valgrind

→ `valgrind ./a.out 3`

- Sanitizer Address (ASAN)

→ `clang++ -fsanitize=address -fno-omit-frame-pointer -g 1-OutOfBounds.cpp`

Demo

I-Out of bounds

- Summary :
 - Valgrind
 - Old tool (>15 years). Works with C++, Fortran, etc...
 - Find out of bounds for heap, but not for stack !
 - Can be slow (10x, 50x slower than original code).
 - Sanitizer Address (ASAN)
 - More recent tool. C++ specific.
 - Find more cases (stack and heap)
 - Much faster (2x, 3x slower than original code).

Demo

2-Invalid-iterator

- A Typical error in C++. Pointer or iterator invalidation.

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v(10, 3);           // 10 numbers init to 3
    auto p = v.begin();                 // an iterator on the beginning of the vector
    std::cout << *p << std::endl;      // print the first value : should be 3
    v.push_back(7);                     // Add one more number to the vector,
    std::cout << *p << std::endl;      // print again the first value ???
}
```


2-Invalid-iterator

- A Typical error in C++. Pointer or iterator invalidation.

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v(10, 3);           // 10 numbers init to 3
    auto p = v.begin();                 // an iterator on the beginning of the vector
    std::cout << *p << std::endl;      // print the first value : should be 3
    v.push_back(7);                     // Add one more number to the vector,
    std::cout << *p << std::endl;      // print again the first value ???
}
```

- Next generation of compiler will detect this.
Lifetime Proposal for std C++.
- Experimental branch of clang.

→ clang++ -Wlifetime 7-IteratorInvalid.cpp

<https://godbolt.org/z/Z878Dp>

3- Memory leak


- Allocate some memory, and ... forget to release it

```
int main() {  
    int *g = new int[100]; // allocate the array of 100 int on the heap  
    g = nullptr;           // Lost the pointer.  
  
    // ... whatever ...  
}
```

- **Effect** : Often none, except if you call such a function a lot, you will run out of memory !
- **Modern C++**: no new/delete. Should not pass code review !
- **Tools** : Valgrind, ASAN (address sanitizer).
 - `valgrind --leak-check=full ./a.out`
 - `clang++ -fsanitize=address -g Leak.cpp`

4-Uninitialized variables

- What happens if we forgot to initialise something ?



```
#include <iostream>

int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int
    int factorial; // OOPS !
    for (int i = 1; i <= num; ++i) {
        factorial *= i;
    }
    std::cout << num << "! = " << factorial << "\n";
    return factorial;
}
```

- Static analyser
 - Compiler (clang -Wall detects it, but not gcc)
 - Cppcheck
- Dynamical analyzer: Valgrind, memory sanitiser (MSAN)
- **Limitation** : libraries must be compiled with `-fsanitize=memory`

Demo

→ `clang++ -g -fsanitize=memory -fno-omit-frame-pointer -fsanitize-memory-track-origins factorial.cpp`

5- Undefined behaviour

- Undefined behavior sanitizer finds many things :
 - Overflow
 - Division by zero
 - Wrong cast e.g. in calling C lib
 - ... Undefined Behavior situation in C++...

5- Undefined behaviour

- Integer overflow.
- << is binary shift, we just shifted too much !

```
#include <iostream>
int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int

    int t = num << 16;
    int r = t * t;
    std::cout << r << std::endl;

    double x = 1/0.0;
    std::cout << x << std::endl;
}
```

```
→ clang++ -Wall int_overflow.cpp
→ ./a.out 5
0
→ clang++ -Wall -fsanitize=undefined int_overflow.cpp
→ ./a.out 5
int_overflow.cpp:7:18:
runtime error: signed integer overflow: 327680 * 327680 cannot be represented in type 'int'
0
```

6-Logic error

- The logic of the code is flawed, but no automatic tool will find this !

```
#include <iostream>

int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int

    int factorial = 0; // Oops !
    for (int i = 1; i <= num; ++i) {
        factorial *= i;
    }
    std::cout << num << "! = " << factorial << "\n";
    return factorial;
}
```

Demo

- Let's follow step by step and use the debugger.
- GDB : for many compiled languages: C, C++, FORTRAN, Java, Pascal, Ada, D, Go
- Gdbgui : a light interface to GDB (usable remotely). Or use an IDE

7-Infinite Loop

- Due to some logic flaw, your code is stuck in an infinite loop !

```
int main() {  
  
    int c = 0;  
  
    while (1) {  
        c += 1;  
        // whatever  
    }  
}
```

- You can run, and attach gdb to your process on the fly

```
→ gdb ./a.out --pid=PID_OF_THE_PROCESS
```

Demo

GDB cheatsheet - page 1

Running

```
# gdb <program> [core dump]
    Start GDB (with optional core dump).

# gdb --args <program> <args...>
    Start GDB and pass arguments

# gdb --pid <pid>
    Start GDB and attach to process.

set args <args...>
    Set arguments to pass to program to
    be debugged.

run
    Run the program to be debugged.

kill
    Kill the running program.
```

Breakpoints

```
break <where>
    Set a new breakpoint.

delete <breakpoint#>
    Remove a breakpoint.

clear
    Delete all breakpoints.

enable <breakpoint#>
    Enable a disabled breakpoint.

disable <breakpoint#>
    Disable a breakpoint.
```

Watchpoints

```
watch <where>
    Set a new watchpoint.

delete/enable/disable <watchpoint#>
    Like breakpoints.
```

<where>

```
function_name
    Break/watch the named function.

line_number
    Break/watch the line number in the cur-
    rent source file.

file:line_number
    Break/watch the line number in the
    named source file.
```

Conditions

```
break/watch <where> if <condition>
    Break/watch at the given location if the
    condition is met.
    Conditions may be almost any C ex-
    pression that evaluate to true or false.

condition <breakpoint#> <condition>
    Set/change the condition of an existing
    break- or watchpoint.
```

Examining the stack

```
backtrace
where
    Show call stack.

backtrace full
where full
    Show call stack, also print the local va-
    riables in each frame.

frame <frame#>
    Select the stack frame to operate on.
```

Stepping

```
step
    Go to next instruction (source line), di-
    ving into function.
```

```
next
    Go to next instruction (source line) but
    don't dive into functions.

finish
    Continue until the current function re-
    turns.

continue
    Continue normal execution.
```

Variables and memory

```
print/format <what>
    Print content of variable/memory locati-
    on/register.

display/format <what>
    Like „print“, but print the information
    after each stepping instruction.

undisplay <display#>
    Remove the „display“ with the given
    number.

enable display <display#>
disable display <display#>
    En- or disable the „display“ with the gi-
    ven number.

x/nfu <address>
    Print memory.
    n: How many units to print (default 1).
    f: Format character (like „print“).
    u: Unit.

    Unit is one of:
        b: Byte,
        h: Half-word (two bytes)
        w: Word (four bytes)
        g: Giant word (eight bytes)).
```


GDB cheatsheet - page 2

Format

<i>a</i>	Pointer.
<i>c</i>	Read as integer, print as character.
<i>d</i>	Integer, signed decimal.
<i>f</i>	Floating point number.
<i>o</i>	Integer, print as octal.
<i>s</i>	Try to treat as C string.
<i>t</i>	Integer, print as binary (<i>t</i> = „two“).
<i>u</i>	Integer, unsigned decimal.
<i>x</i>	Integer, print as hexadecimal.

<what>

expression

Almost any C expression, including function calls (must be prefixed with a cast to tell GDB the return value type).

file_name::variable_name

Content of the variable defined in the named file (static variables).

function::variable_name

Content of the variable defined in the named function (if on the stack).

{*type*}*address*

Content at *address*, interpreted as being of the C type *type*.

\$*register*

Content of named register. Interesting registers are \$esp (stack pointer), \$ebp (frame pointer) and \$eip (instruction pointer).

Threads

thread <*thread#*>

Chose thread to operate on.

Manipulating the program

set var <*variable_name*>=<*value*>

Change the content of a variable to the given value.

return <*expression*>

Force the current function to return immediately, passing the given value.

Sources

directory <*directory*>

Add *directory* to the list of directories that is searched for sources.

list

list <*filename*>:<*function*>list <*filename*>:<*line_number*>list <*first*>,<*last*>

Shows the current or given source context. The *filename* may be omitted. If *last* is omitted the context starting at *start* is printed instead of centered around it.

set listsize <*count*>

Set how many lines to show in „list“.

Signals

handle <*signal*> <*options*>

Set how to handle signles. Options are:

(*no*)*print*: (Don't) print a message when signals occurs.

(*no*)*stop*: (Don't) stop the program when signals occurs.

(*no*)*pass*: (Don't) pass the signal to the program.

Informations

disassemble

disassemble <*where*>

Disassemble the current function or given location.

info args

Print the arguments to the function of the current stack frame.

info breakpoints

Print informations about the break- and watchpoints.

info display

Print informations about the „displays“.

info locals

Print the local variables in the currently selected stack frame.

info sharedlibrary

List loaded shared libraries.

info signals

List all signals and how they are currently handled.

info threads

List all threads.

show directories

Print all directories in which GDB searches for source files.

show listsize

Print how many are shown in the „list“ command.

whatis *variable_name*

Print type of named variable.

8-Thread/OpenMP

- Race condition. Can you spot it/them ?

```
#include <omp.h>
#include <iostream>
#include <vector>

int main(int argc, char* argv[]) {
    std::vector<double> data(100, 1.0);
    double sum;

    #pragma omp parallel shared(sum, data)
    {
        sum = 0.0;
        #pragma omp for
        for(int i = 0; i < data.size(); i++){
            sum += data[i];
        }
    }
    std::cout << sum << std::endl;
}
```

→ clang++ -fopenmp -g -fsanitize=thread 8-RaceCondition.cpp

- Experimental branch of clang.

Good practices

- Clear, expressive code (Modern C++)
- Code Review !
- **Automated Tests** (googletest, TDD)
- Version Control (git)
- **Static analyzer** (clang-tidy)
- **Dynamic analyzer** (Valgrind, clang sanitizer)
- Logic error. Debugger, step by step. **GDB**
- Print stuff ? Only in the last resort ...

Thank you for your attention!